2D MPI + CUDA | n=2048 | one versus two GPUs

Vanda | audited GPU measurements | 3 independent runs per configuration

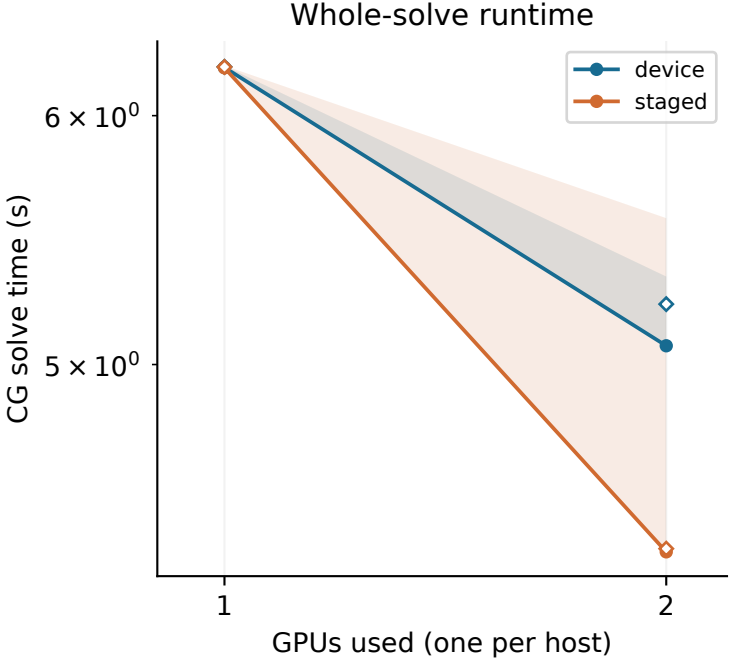

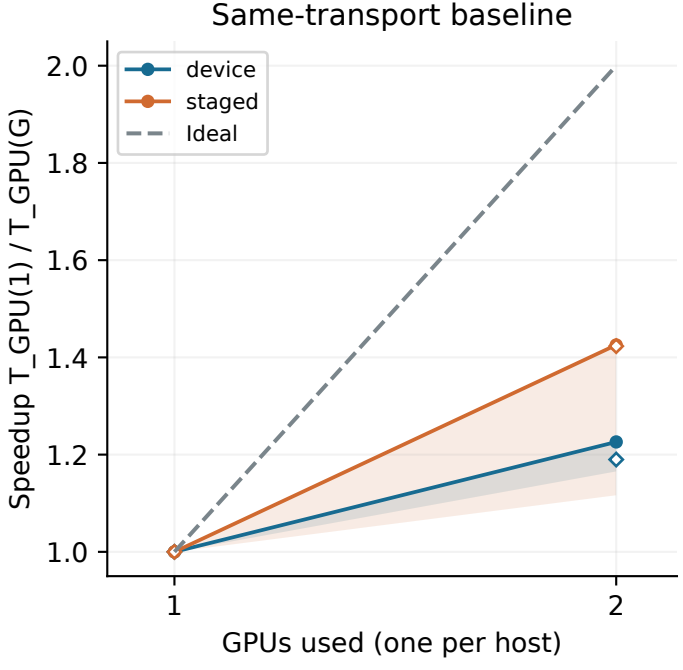

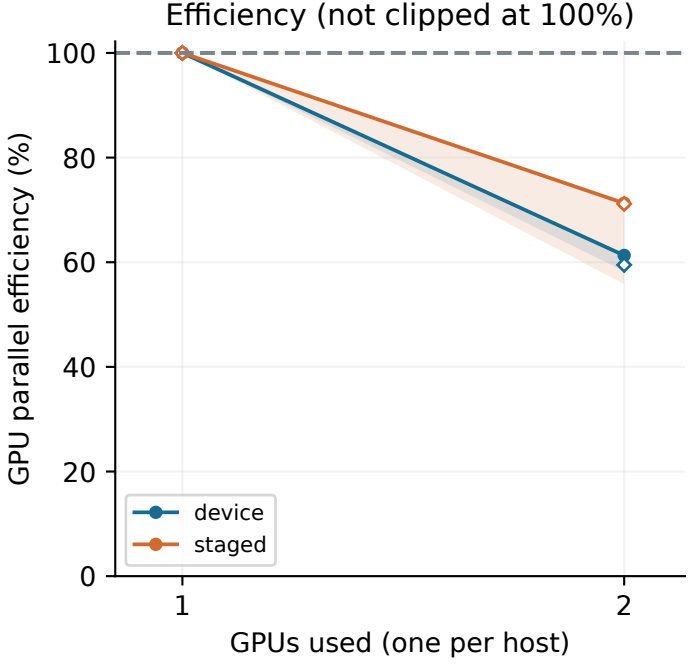

Line/filled marker: minimum; hollow diamond: median; shading: observed min-max (3 runs, not a CI).
Each transport uses its own one-GPU baseline in the same allocation. Device-pointer MPI over TCP may internally stage.